

Efficient data transmission with TCP

In this lesson, we'll study the main data transfer mechanisms used by TCP.

We'll cover the following

- Segment Transmission Strategies
- Nagle's Algorithm
 - Algorithm
 - Limitations
- Quick Quiz!

Segment Transmission Strategies

In a transport protocol such as TCP that offers a byte stream, a practical issue that was left as an implementation choice in [RFC 793](#) was *when* a new TCP segment should be sent.

There are two simple and extreme implementation choices:

1. Send a TCP segment as soon as the application has requested the transmission of some data.
 - **Advantage:** This allows TCP to provide a **low delay service**.
 - **Disadvantage:** If the application is writing data one byte at a time, TCP would place each byte in a segment containing 20 bytes of the TCP header. This is a **huge overhead** that is not acceptable in wide area networks.
2. Transmit a new TCP segment once the application has produced MSS bytes of data. Recall MSS from this lesson on [TCP Headers](#).
 - **Advantage: Reduced overhead**
 - **Disadvantage:** Potentially at the cost of a **very high delay**, which may be unacceptable for interactive applications.

Nagle's Algorithm

An elegant solution to this problem was proposed by John Nagle in [RFC 896](#) called **Nagle's Algorithm**.

In essence, as long as there are unacknowledged packets, Nagle's algorithm keeps collecting application layer data up to maximum segment size, to be sent together in a single packet. This helps reduce the packetization overhead by reducing small packets.

Algorithm

```
if window size >= MSS and available data >= MSS:
    send one MSS-sized segment
else:
    if there is unacknowledged data:
        place data in buffer until acknowledgement is received
    else:
        send one TCP segment containing all buffered data
```

Nagle's Algorithm

The algorithm takes up 2-3 lines in a TCP implementation. Nagle's algorithm 'executes' every time new data comes in from the remote host. Here's how it works: it sends data if it is at least the size of one MSS and the window size is appropriate. Otherwise, it checks if any unacknowledged segments exist. If so, it **buffers the data** and doesn't send it. There is no timer on this condition, and it will keep buffering data until previous segments are acknowledged. If an *ACK* segment comes in, it sends the data.

Limitations

Nagle's has a few limitations:

1. Nagle's algorithm is only supported by TCP and no other protocols like UDP.
2. TCP applications that require **low latency** and **fast response times** such as internet phone calls or real-time online video games, do not work well when Nagle's is enabled. The delay caused by the algorithm triggers a noticeable lag. These applications usually *disable* Nagle's with an interface called the `TCP_NODELAY` option.

3. The algorithm was originally developed at a time when computer networks supported much less bandwidth than they do today. It saved bandwidth and made a lot of sense at the time, however, the algorithm is much less frequently used today.
4. The algorithm also works poorly with **delayed ACKS**, a TCP feature that is used now. With both algorithms enabled, applications experience a consistent delay because Nagle's algorithm doesn't send data until an ACK is received and delayed ACKs feature doesn't send an ACK until after a certain delay.

Quick Quiz!

1

Consider the following scenario:

- The last segment received by a TCP entity had an ACK value of 10.
- The MSS value is 536.
- There are 50 bytes in the buffer waiting for transmission.
- The window size is 500.
- The application sends a 500 byte message.
- All the data sent so far has been acknowledged.

Will the data be buffered or sent on if Nagle's was active?

- ☐ A) It will be buffered
- ☐ B) One TCP segment with all the buffered data will be sent
- ☐ C) One MSS sized segment will be sent

For now, that's all on Nagle's. Let's move on to TCP window-scaling mechanisms.