

# TCP Congestion Control: AIMD

The last part of TCP that we're going to study is congestion control.

## We'll cover the following

- Requirements of a Congestion Control Algorithm
- TCP Congestion Control Algorithms
  - Additive Increase Multiplicative Decrease
- Quick Quiz!

We've already looked at what congestion control is in a [previous lesson](#). Let's get into some TCP-specific congestion control algorithms now!

But first:

## Requirements of a Congestion Control Algorithm

#

1. The congestion control scheme **must avoid congestion**. In practice, this means that the algorithm should ensure that the bandwidth allocated to a certain host at any given time does not exceed the bandwidth of the bottleneck link.
2. The congestion control scheme **must be efficient**. The congestion control scheme should ensure that the available bandwidth is efficiently used. Namely that the bandwidth allocated to a certain host at any given time doesn't fall too far below the bandwidth of the bottleneck link.
3. The congestion control scheme **should be fair**. Most congestion schemes aim at achieving max-min fairness, which we had a lesson dedicated to.

## TCP Congestion Control Algorithms #

To implement most TCP congestion control algorithms, a TCP host must be able to control its **transmission rate**. In order to do so, it can constrain its

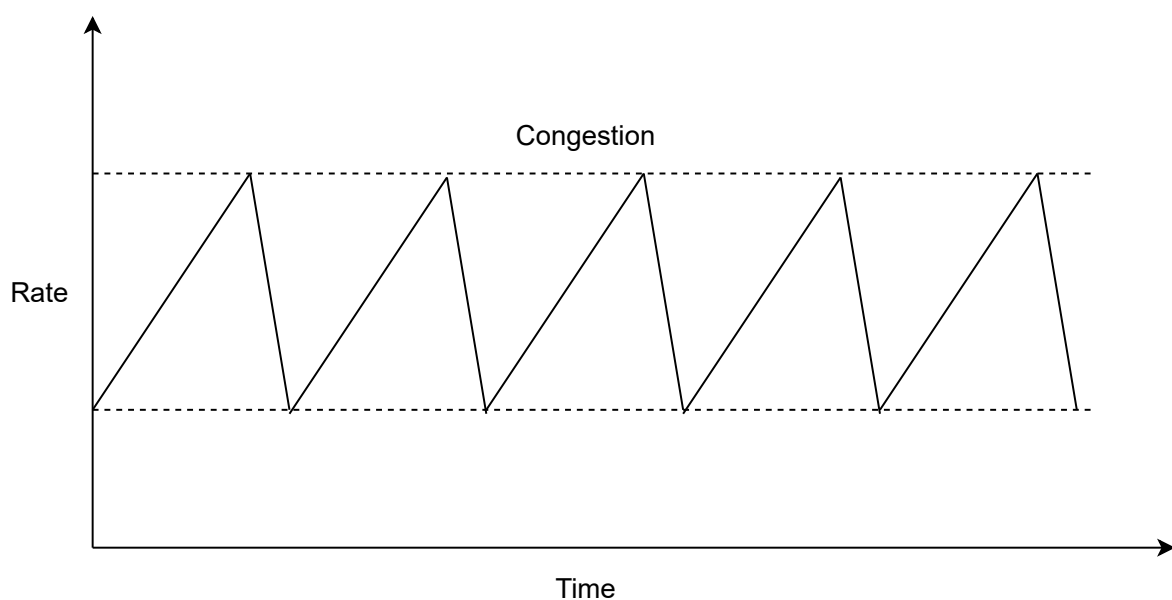
**sending window.** Recall that after transmitting data equal to the window size, the sender must pause and wait at least one RTT for the *ACK* before it can transmit more data. Thus the maximum data rate is:  $\frac{window}{rtt}$  where *window* is the maximum between the host's sending window and the window advertised by the receiver.

Hence, TCP's congestion control scheme can restrict sending windows based on a **congestion window**. The current value of the congestion window is stored in the TCB of each TCP connection. The value of the window that can be used by the sender is  $\min(\text{congestion window}, \text{receiving window}, \text{sending window})$ .

## Additive Increase Multiplicative Decrease #

The **Additive Increase Multiplicative Decrease** algorithm *decreases* the transmission rate of a host when congestion has been detected in the network. Congestion is detected when acknowledgments for packets do not arrive before the transmission timer times out. Furthermore, it *increases* their transmission rate when the network is *not* congested. Hence, the rate allocated to each host fluctuates with time, depending on the feedback received from the network.

The figure below illustrates the evolution of the transmission rates allocated to a host in a network.



How the transmission rate of a host changes with AIMD. Notice the 'saw tooth' pattern.

The **Additive Increase** part of the TCP congestion control increments the congestion window by *MSS* bytes every round-trip time. In the TCP literature, this phase is often called the *congestion avoidance* phase. Once congestion is detected, the **Multiplicative Decrease** part of the TCP congestion control reacts by multiplying the current value of the congestion window with a number greater than 0 and less than 1.

However, since the increase in the window is so slow, the TCP connection may have to wait for many round-trip times before being able to efficiently use the available bandwidth. To avoid this, the TCP congestion control scheme includes the **slow-start** algorithm.

## Quick Quiz! #

1

Which of the following is not a requirement of a congestion control algorithm?

- ☐ A) Must avoid congestion
- ☐ B) Must be fair
- ☐ C) Must be efficient
- ☐ D) Must be accurate

COMPLETED 0%

1 of 2



Let's have a look at the slow start algorithm in the next lesson.

